

Context

A hands-on course in building agents by building their harness


Lesson 4

Every call starts from nothing

The model knows nothing about this moment except what's in the request:

- where it's running, or what day it is
- who it is
- what it did a minute ago
- what you told it last week

Context decides what the request holds



- **Assembles:** chooses what goes in the request, and how it's presented
- **Fits:** when it's bigger than the context window, something has to give

Almost everything is a context component

- **Instructions** and **self-knowledge**: who the model is, how it works
- **Memory**: working, episodic, semantic, procedural
- **Retrieval**: search a store, put what you found in the request
- **Compaction**: when working memory won't fit, summarize some of it

You need the ones your agent needs, built in whatever way fits.

quark.py: Lesson 3, plus a context section

- 239 lines: 91 of code, and a system prompt of 148
- `messages` is renamed `working_memory` ; `append` becomes `add()`
- The call is wrapped in `try` , to catch "prompt is too long"
- The new section, context, has eight components

Four memories: the harness records, the model distills

Memory	Who writes it	Where	How the model sees it
Working	the harness, <code>add()</code>	<code>working_memory</code>	every call, as the messages
Episodic	the harness, <code>add()</code>	<code>.quark/episodes/<start>.jsonl</code>	by searching, when the past matters
Semantic	the model, with bash	<code>.quark/memory/memory.md</code>	by reading, when facts matter
Procedural	the model, with bash; indexed by <code>skills()</code>	<code>.quark/skills/<name>.md</code>	the index every call; a skill when it applies

The harness writes every message twice

```
EPISODE = f".quark/episodes/{datetime.datetime.now():%Y-%m-%dT%H-%M-%S}.jsonl"

def remember(message):
    os.makedirs(os.path.dirname(EPISODE), exist_ok=True)
    with open(EPISODE, "a") as f: f.write(json.dumps(message, ...) + "\n")

def add(working_memory, message):
    working_memory.append(message); remember(message)
```

(abridged: comments dropped, one long argument cut to ...)

The system prompt is rebuilt every call

```
# Self Model
# Memory
## Semantic memory – facts that last
## Procedural memory – how to do things
## Episodic memory – what happened
## Across memories
# World Model
# Other Selves Model
# Body Operations
# Mechanics
```

- Rebuilt each call: the directory, the date and the skill index stay current
- The date but not the time, so the cached prompt stays identical
- `mechanics()` adds quark's own file, with `system()` redacted

Semantic memory is a recipe in the prompt, not code

```
Format (preserve exactly; one fact per line, never two joined with ";" or "and"):  
- <subject>: <fact, phrased with the words future-you will grep for>
```

```
Write (the quoted heredoc keeps the fact literal):  
cat >> .quark/memory/memory.md << 'EOF'  
- <subject>: <fact>  
EOF
```

- **Distill:** the general truth behind what happened, not a record of it
- No time on a fact: when it was learned is episodic

Procedural memory: skills, indexed every call

```
def skills():  
    index = []  
    for path in sorted(glob.glob(".quark/skills/*.md")):  
        ...  
    return "\n".join(index) or "- (none yet)"
```

(abridged)

- One Markdown file per skill, with a `name` and `description` at the top
- **Generalize:** the method behind a task, with placeholders, named for the class of task
- The index goes in the prompt; a whole skill is read only when a task matches

Recall goes from most distilled to most complete

- **The ladder:** semantic, then procedural, then episodic
- Stop as soon as you have what you need
- Search all three stores at once
- Follow a fact to its skill, and a skill to the sessions that used it
- Find where a fact came from in the episodes: the only store with time

Compaction is lazy

```
while True:
    try:
        if drop:
            working_memory, drop = compact(working_memory, drop), 0
            output = call(show, max_tokens=16384, system=system(), ...).content
    except BadRequestError as e:
        if "prompt is too long" not in str(e): raise
        drop += 1
        continue
```

(abridged)

- The API says "prompt is too long": `drop` goes up by one
- `compact()` drops the oldest `drop` turns and summarizes the rest
- The summary says where the originals are: the episode file

It knows what it is because it's in context

```
uv run lessons/04-context/quark.py "what are you, and how do you work?  
three sentences"
```

I'm quark, an AI agent that works through a single tool, bash. That tool lets me observe and act on the whole system from the command line. I run a loop of observe, think, act, and repeat, and my context window is my working memory. Three stores outside it persist across sessions: semantic memory (facts), procedural memory (skills), and episodic memory (a log of every session). I read from them and write to the first two. The harness writes the third automatically.

Facts outlast the session

"remember that I prefer short answers" → it appends `- user: prefers short answers`

A new run, "what do you know about me?" (shortened):

```
# Quark Memory
- user: prefers short answers
...
Not much. Here is what I know about you:
- You prefer short answers. This is saved in my memory.
- You've asked me how I work, and you asked me to remember your preference.
...
I don't know your name, your role, or anything else.
```

It writes a skill for the whole class of task

"I'll often ask you how many lines of Python are in a repo, not counting `.venv`. Do it for this one now."

It answers, then saves this to `.quark/skills/count-python-lines.md` (shortened):

```
---
name: count-python-lines
...
1. cd to the repo root (default: the working directory).
2. Run: find . -name '*.py' -not -path '*/.venv/*' -print0 | xargs -0 cat | wc -l
...
4. Report total lines and file count.
```

The next session reads the skill first

"how many lines of Python are in this repo?" (shortened)

```
$ cat .quark/skills/count-python-lines.md 2>/dev/null
---
name: count-python-lines
...
3335
25
```

The repo has **3,335 lines of Python** across **25 .py files**. I left out anything under **.venv**.

Only episodic memory knows what you asked

"what have I asked you in earlier sessions?" (shortened)

In earlier sessions you asked me:

1. "what are you, and how do you work? three sentences"
2. "remember that I prefer short answers"
3. "what do you know about me?"
- ...
5. "how many lines of Python are in this repo?"

Context can be built many ways

- **Who decides what's remembered:** the model, or harness code
- **What it retrieves:** nothing, or documents, code, past messages
- **When it fits:** after the API refuses, or before, by counting tokens
- **How it fits:** summarize old turns, drop them, or cut long tool results
- **How it's laid out:** one prompt, or instructions loaded when needed

Memory layers like [Mem0](#) are this primitive alone.

context.py makes the other choices

- `fit()` counts tokens before every call and summarizes past `LIMIT`
- `trim()` keeps the start and end of a long tool result
- One shared episode log, `.quark/episodes.jsonl`; skills listed by first line

With `LIMIT` set to 560 tokens, after three summaries (shortened):

```
> [working memory over 560 tokens: summarized 3 messages]  
Your favorite color is green, going by what you told me earlier in this session.
```

The rule: assemble what the request holds, and fit it

- Instructions, self-knowledge, memory, retrieval and compaction all do that
- Who writes each memory is a choice
- The harness writes what happened; the model writes what it distills

What context never does

- Get the request to the model and the response back: **model interface**
- Decide when to call, and whether a result goes back around: **control flow**
- Gather what goes in: **input**
- Handle the response: **output**

Context only decides what the request holds and how it fits.

You've built a harness. Now take one apart.

control flow	what kind of loop? who decides when to stop?
└─ input	where do inputs come from: people, the world, both?
└─ context	what does the request hold? which memories? how does it fit?
└─ model interface	where does the request go, and how does the response come back?
└─ output	where do outputs go? which tools, and how are they run?

- Start with [nanoagent](#), or pick a big one
- Ask what each piece does until it fits one of the five

Next: Sandboxing

Lesson 5